

Flyby-F11 Autonomy Architecture: Ontology-Constrained Reinforcement Learning

Finley Holt

2025-12-26

Table of contents

1 Flyby-F11 Autonomy Architecture: Ontology-Constrained Reinforcement Learning	2
1.1 Executive Summary	2
1.1.1 Why This Approach?	2
1.2 Architecture Overview	2
1.2.1 Single-Reasoner Architecture (Vampire Only)	2
1.2.2 Tiered Safety Architecture	3
1.2.3 Unified Compute Architecture	3
1.2.4 Three-Level Hierarchy (Execution Mode)	4
1.2.5 Integration Points	5
1.3 Component Details	5
1.3.1 1. Ontological Knowledge Base	5
1.3.2 2. Multi-Agent Reinforcement Learning	7
1.3.3 3. Automatic Goal Generation Model (AGGM)	9
1.3.4 4. Safety Constraint Enforcement	10
1.3.5 5. Perception-to-Reasoning Bridge	11
1.4 Implementation Roadmap	15
1.4.1 Phase 1: Ontology Foundation & Two-Phase Architecture (Weeks 1-6)	15
1.4.2 Phase 2: Perception-to-Reasoning Bridge (Weeks 7-10)	16
1.4.3 Phase 3: Multi-Agent RL (Weeks 9-14)	16
1.4.4 Phase 4: Benchmark Evaluation (Weeks 15-18)	17
1.4.5 Phase 5: Hardware Validation (Weeks 19-24)	17
1.4.6 Phase 6: Documentation (Weeks 25-26)	17
1.5 Key Innovations and Contributions	17
1.5.1 1. Two-Phase Compute Architecture for Edge Ontological Reasoning	17
1.5.2 2. Perception-to-Reasoning Bridge for Symbolic Grounding	18
1.5.3 3. Real-Time Ontological Reasoning on Edge Hardware	18
1.5.4 4. AGGM for GPS-Denied Navigation	18
1.5.5 5. Explainable Autonomy for Defense Applications	18
1.5.6 6. Hierarchical Multi-Agent RL with Ontological Structure	18
1.6 Success Criteria	19
1.6.1 Simulation Benchmarks (Phase 4)	19
1.6.2 Hardware Performance (Phase 5)	19

1.6.3	Flight Testing (Phase 5)	19
1.6.4	Research Impact (Phase 6)	19
1.7	References	19
1.8	Appendix: Alignment with Project-Drone	20

1 Flyby-F11 Autonomy Architecture: Ontology-Constrained Reinforcement Learning

Platform: Flyby Robotics F-11 Developer Quadcopter with Jetson Orin NX 16GB

Context: GPS-denied, communications-limited autonomous missions (MCTSSA collaboration)

Core Innovation: Hybrid ontology-constrained RL for safe, adaptive, explainable autonomy

1.1 Executive Summary

This document defines the architectural approach for autonomous navigation on the Flyby-F11 platform. Based on comprehensive literature review ([SYNTHESIS.qmd](#)), we implement a **hybrid architecture** combining:

1. **Formal Ontological Knowledge** - SUMO-based ontology for safety constraints, domain vocabulary, and semantic reasoning
2. **Multi-Agent Reinforcement Learning** - Hierarchical RL agents optimizing within ontology-defined safe action spaces
3. **Automatic Goal Generation Model (AGGM)** - Runtime adaptation to unseen situations through ontological reasoning

This approach addresses critical failures in pure learning systems (37-65% collision rates in UAV-ON benchmark) while maintaining adaptability that pure rule-based systems lack.

1.1.1 Why This Approach?

Evidence from Literature ([detailed synthesis](#)): - **Safety:** Ontologies reduce collision rates by constraining RL action spaces before execution - **Adaptability:** AGGM enables runtime goal generation for unseen situations without retraining - **Explainability:** Semantic rule tracing provides audit trails required for defense applications (MCTSSA) - **Efficiency:** Ontological abstractions improve sim-to-real transfer and sample efficiency - **Edge-Compatible:** <100ms reasoning latency achievable on Jetson platforms

1.2 Architecture Overview

1.2.1 Single-Reasoner Architecture (Vampire Only)

Architectural Decision (Phase 3 Evaluation, 2024-12-25)

After comprehensive empirical benchmarking ([Phase 3 Evaluation Report](#)), we determined that a **single-reasoner architecture using Vampire** is optimal:

Finding	Impact
Vampire ~50ms p95 latency	Acceptable for 20Hz navigation loop
OWL reasoners cannot express safety axioms	ELK/Reasonable rejected (DL limitations)
Prolog 4,700x faster but adds complexity	Rejected to avoid translation/maintenance burden
KIF/SUMO remains single source of truth	0% translation loss

Key Insight: Ontological reasoning belongs in the navigation layer (20Hz), not the flight control layer (400Hz). Real-time safety (<10ms) is handled by the classical control layer, not symbolic reasoning.

1.2.2 Tiered Safety Architecture

TIER 1: Classical Control (<1ms)

PX4/ArduPilot attitude control

Motor mixing, PID loops

Sensor filtering, state estimation

TIER 2: Pre-computed Safety (<10ms)

Obstacle costmaps (baked from depth)

Geofence boundary polygons

Velocity limits, acceleration constraints

TIER 3: Tactical Reasoning (~50ms) - VAMPIRE

"Am I violating a no-fly zone?"

"Is battery critical for return?"

"Does this waypoint sequence satisfy constraints?"

Runs at 20Hz navigation rate

TIER 4: Mission Planning (~100ms-1s) - VAMPIRE

Route verification and regulatory compliance

Mission feasibility analysis

Pre-flight or during mission replanning

1.2.3 Unified Compute Architecture

With the single-reasoner decision, there is no mode swapping between planning and execution:

UNIFIED EXECUTION MODE

Vision/Perception (~12-14 GB)	Vampire Reasoning (~50-100 MB)
----------------------------------	-----------------------------------

YOLO11 (TensorRT)	Vampire Theorem Prover
~2-3 GB	~14 MB (typical)
- Object detection	- Full FOL reasoning
- 20-50ms latency	- Safety queries: ~50ms
Segmentation Model	- Planning queries: ~100ms
~1-2 GB	- Query caching for repeated
- Terrain types	Perception → TPTP Bridge
- Traversability	~50 MB
VLM (Optional)	- Vision facts → TPTP format
~3-5 GB (7B model)	- Spatial relations computed
- Scene understanding	- Event detection
- Semantic grounding	ROS 2 vampire_bridge Package
Depth Processing	- Query service (100ms timeout)
~500 MB	- Result caching (LRU + TTL)
- Obstacle maps	- 20Hz safety monitoring
- Clearance calc	
ROS 2 Middleware + System Overhead: ~2 GB	

Key Advantages: - **No translation needed:** KIF/SUMO is the single source of truth - **Full expressivity:** All safety axioms preserved (0% loss) - **Simplified architecture:** One reasoner for planning AND runtime - **Acceptable latency:** ~50ms fits within 20Hz navigation loop

1.2.4 Three-Level Hierarchy (Execution Mode)

LEVEL 1: MISSION PLANNING (10-second horizon)

Vampire (TPTP)	RL Agent: Mission
- Mission rules	- Select waypoints
- Constraints	- Adapt mission
- Airspace rules	- Resource planning

Waypoint sequence

LEVEL 2: BEHAVIOR SELECTION (1-second horizon)

Vampire (TPTP)	RL Agent: Behavior
- Behavior rules	- Navigate
- Transitions	- Loiter
- Safety rules	- Land

- Avoid

Behavior + parameters

LEVEL 3: TRAJECTORY OPTIMIZATION (100-ms horizon)

Vampire (TPTP)	RL Agent: Trajectory
- Actuator limits	- Velocity commands
- Safety margins	- Obstacle avoid
- Clearances	- Energy optimize

[vx, vy, vz, yaw]

PX4/ArduPilot
(Control Loop)

1.2.5 Integration Points

Single Reasoner (No Mode Swapping): - Vampire handles both planning and runtime tactical queries - Different query types distinguished by timeout (planning: 5s, tactical: 100ms) - Query caching reduces repeated query overhead (LRU + TTL) - KIF/SUMO remains single source of truth with 0% translation loss

Ontology → RL (Constraint Enforcement): - Filters action spaces: Vampire proves which actions satisfy constraints - Shapes rewards: penalties for ontology-detected violations - Structures state: semantic abstractions from TPTP facts

RL → Ontology (Experience-Based Learning): - Optimizes within safe boundaries defined by ontology - Discovers efficient policies through exploration - Adapts parameters based on environmental feedback

Bidirectional (AGGM Runtime Reasoning): - Forward reasoning: select goals from ontology-defined goal set - Backward reasoning: create new goals to return to known safe states - Importance weighting: prioritize safety-critical observations

1.3 Component Details

1.3.1 1. Ontological Knowledge Base

Foundation: SUMO upper-level ontology (design/verification) + SWI-Prolog (runtime inference)
Reasoning Strategy: Two-phase architecture detailed in [ONTOLOGY_FOUNDATION.qmd](#)

1.3.1.1 Planning Phase: SUMO + Heavyweight Reasoners

SUMO Ontology (SUO-KIF format): - Full first-order logic with n-ary relations - ~25,000 terms, ~80,000 axioms - Supports ternary+ relations: (orientation ?OBJ1 ?OBJ2 ?DIRECTION) - Used for mission modeling, verification, safety proofs

Reasoning Engines: - **Vampire:** Automated theorem proving for safety property verification - **Clingo:** Answer set programming for optimal path planning - **E-Prover:** Alternative FOL reasoner for constraint checking

Four-Layer Structure:

Upper Level: SUMO

- Physical objects (Object, Agent)
- Processes (Motion, Action)
- Relations (orientation, between, during)
- Functions (distance, measure, direction)

Domain Level: IEEE AUR + UAV Extensions

- UAV (extends TransportationDevice)
- FlightPhase (Takeoff, Transit, Landing)
- SpatialRelation (hasSafeSeparation)
- EnvironmentalCondition (Wind, Visibility)

Application Level: Flyby-F11 Missions

- SurveyMission, InspectionMission
- GeofenceBoundary, NoFlyZone
- ISR payloads (Gremsy VIO, RESEPI, Raptor)
- NDAACompliance, FAAPart107Rules

Instance Level: Runtime State

- currentMission: Inspection123
- detectedObstacle: Tree47 at [x,y,z]
- batteryLevel: 68% (12 min remaining)

1.3.1.2 Execution Phase: Compiled Prolog Rules

SWI-Prolog Runtime (~50-100 MB): - Compiled from SUMO axioms relevant to mission - Optimized for <10ms query latency - Engine architecture: ~20KB per concurrent query thread - TCMalloc for reduced memory footprint

Example Compiled Rules (Prolog syntax):

```
% Collision avoidance (from SUMO spatial reasoning)
canExecute(moveToward(Pos)) :-
    forall(obstacle(Obs),
        distance(currentPosition, Obs, Dist),
        Dist > safetyMargin).

% Energy management (from SUMO resource ontology)
mustReturnToHome :-
    estimatedEnergyRemaining(Energy),
    energyToReturnHome(Required),
    safetyReserve(Reserve),
    Energy < (Required + Reserve).

% Geofence enforcement (from SUMO spatial containment)
canExecute(Action) :-
    resultingPosition(Action, Pos),
    isWithinGeofence(Pos).

% Flight phase transitions (from SUMO process ontology)
canTransitionTo(landing) :-
    isLandingZone(currentPosition),
    altitude(Alt), Alt < landingInitiationAltitude,
    horizontalVelocity(Vel), Vel < maxLandingVelocity.

% Spatial relations (ternary - preserved from SUMO)
between(Drone, Obj1, Obj2) :-
    position(Drone, [X1, Y1, Z1]),
    position(Obj1, [X2, Y2, Z2]),
    position(Obj2, [X3, Y3, Z3]),
    % Check if Drone is geometrically between Obj1 and Obj2
    is_on_line_segment([X1,Y1,Z1], [X2,Y2,Z2], [X3,Y3,Z3]).
```

SUMO → Prolog Translation Strategy: 1. **Manual translation** of critical safety axioms (reviewed and verified) 2. **Semi-automatic** for common patterns (spatial relations, temporal logic) 3. **Testing:** Equivalence checking between SUMO proofs and Prolog queries 4. **Validation:** Mission scenarios tested in both planning and execution modes

1.3.2 2. Multi-Agent Reinforcement Learning

Paradigm: Hierarchical RL with experience sharing **Integration:** Each agent operates within ontology-constrained action/state spaces

1.3.2.1 Mission Planner Agent (Level 1)

MDP Formulation: - **State** (ontological): - Mission progress: {waypointsCompleted, currentObjective, remainingObjectives} - Resources: {batteryLevel, timeElapsed, payloadStatus} - Environment: {weatherCondition, airspaceRestrictions, gpssAvailability}

- **Action Space** (ontology-filtered):

- `selectNextWaypoint(waypoint_id)` - only from valid waypoints
- `adaptMissionPlan(new_sequence)` - only safe alternatives
- `abortMission(reason)` - when constraints violated

- **Reward Function:**

```
R = w1 * missionCompletion
  + w2 * efficiencyBonus (time, energy)
  + w3 * safetyMargin (distance to constraints)
  - w4 * constraintViolationPenalty
```

- **Algorithm:** SAC (Soft Actor-Critic) for continuous action parameters
- **Training:** NPS computing cluster with mission scenario diversity

1.3.2.2 Behavior Selector Agent (Level 2)

MDP Formulation: - **State** (ontological): - Vehicle: {altitude, velocity, orientation, flightMode} - Context: {currentWaypoint, obstaclesNearby, terrainType} - Mission: {activeBehavior, missionPhase}

- **Action Space** (ontology-validated):

- `navigate(speed, altitude)` - transit to waypoint
- `loiter(radius, duration)` - station-keeping
- `land(descentRate)` - controlled descent
- `avoid(direction, magnitude)` - collision avoidance maneuver
- `returnToHome()` - emergency return

- **Reward Function:**

```
R = w1 * behaviorAppropriateness (context match)
  + w2 * smoothTransitions (continuity)
  + w3 * constraintSatisfaction (safety)
  - w4 * behaviorSwitchPenalty (stability)
```

- **Algorithm:** PPO (Proximal Policy Optimization) for stable learning
- **Training:** Behavior trees provide structured exploration

1.3.2.3 Trajectory Optimizer Agent (Level 3)

MDP Formulation: - **State** (sensor-based): - Observations: {EKF_pose, ISR_payload_data, YOLO_detections, MAVLink_telemetry} - Dynamics: {position, velocity, acceleration, attitude} - Setpoints: {targetPosition, targetVelocity, targetHeading}

- **Action Space** (actuator-constrained):

- Velocity command: [vx, vy, vz, yaw_rate] within ± 2.5 m/s limits

- **Reward Function:**

```
R = w1 * setpointTracking (error minimization)
  + w2 * trajectorySmooth (jerk minimization)
  + w3 * energyEfficiency (acceleration cost)
```

- + w4 * obstacleClearance (safety margin)
- w5 * collisionPenalty
- **Algorithm:** TD3 (Twin Delayed DDPG) for low-level control
- **Training:** Continuous in simulation with domain randomization

1.3.3 3. Automatic Goal Generation Model (AGGM)

Purpose: Adapt to unseen situations without retraining **Source:** Literature review paper 02 ([Ghanadbashi & Golpayegani, 2022](#))

1.3.3.1 Six-Stage Process

Stage 1: OBSERVE

- Multi-sensor fusion → ontological concepts
- Ontology schema: $L^t = \{C^t, M^t\}$ (concepts, relations)

Stage 2: EVALUATE

- Q-value from RL policy
- State distance: $|s^t - s^{(t-1)}|$
- Importance weighting: $iw_c(x)$ for each concept

Stage 3: IDENTIFY SIGNIFICANT CHANGE (triggers)

- Case 1: Q-value discrepancy (unexpected reward)
- Case 2: State distance threshold (environment change)
- Case 3: High-importance concept detected (safety)

Stage 4: REASON (forward/backward)

- Forward: infer goal from predefined goal-set via SWRL
- Backward: create new goal to reach known safe state
- Priority: $F(B, J)$ balances task vs. state-similarity

Stage 5: GENERATE ACTION

- Ontology-constrained action space (only valid actions)
- RL policy selects action within constraints
- Multi-agent coordination if needed

Stage 6: EXECUTE

- Translate to MAVSDK/MQTT commands
- Monitor execution through telemetry
- Update belief state for next cycle

Example Scenario: Unexpected obstacle during transit

1. **Observe:** ISR payload detects unknown object at 8m distance in flight path
2. **Evaluate:**
 - Q-value drops (expected reward for “move forward” now low)
 - State distance increases (new obstacle concept added)
 - High importance weight (safety-critical observation)
3. **Identify Change:** Case 3 triggered (high-importance safety concept)
4. **Reason:**
 - Forward: No predefined goal matches “unexpected obstacle in path”
 - Backward: Create new goal “reach clear airspace” (return to obstacle-free state)
 - Priority: Safety (J) > task completion (B)
5. **Generate Action:** RL policy selects `avoid(left, 5m)` from ontology-validated actions
6. **Execute:** Send velocity command to PX4, monitor clearance until safe

1.3.4 4. Safety Constraint Enforcement

Critical Requirement: Collision avoidance, geofence compliance, energy management **Method:** Ontology-based action filtering + runtime monitoring

1.3.4.1 Pre-Flight Safety Checks

```
# Ontology evaluates readiness before takeoff
canTakeOff ←
    batteryLevel > minimumForMission
    hasValidLocalization      # GPS or EKF-based position estimate
    ¬hasActiveWarning
    isWithinAuthorizedAirspace
    weatherCondition == Acceptable
```

1.3.4.2 In-Flight Safety Monitoring

```
# Continuous evaluation at 10 Hz
every 100ms:
    # Collision avoidance
    for obstacle in detectedObstacles:
        if distance(uav, obstacle) < safetyMargin:
            trigger_emergency_avoidance(obstacle)

    # Geofence enforcement
    if not isWithinGeofence(current_position):
        trigger_return_to_geofence()

    # Energy management
```

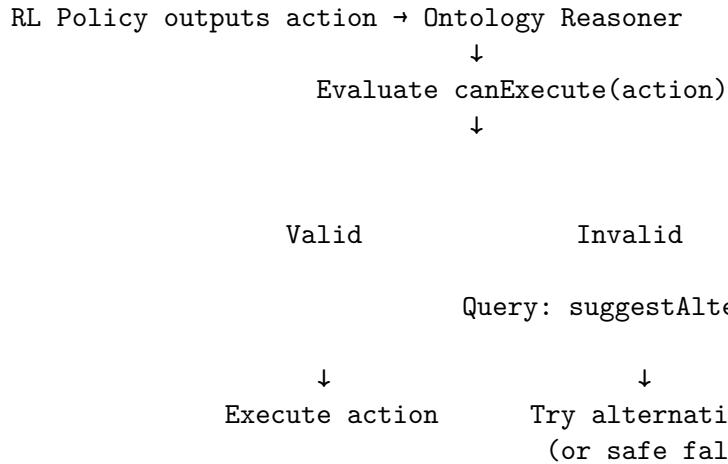
```

if batteryLevel < criticalThreshold:
    trigger_immediate_landing()
if estimatedEnergyRemaining < energyToReturnHome + reserve:
    trigger_return_to_home()

# Sensor health / localization
if not hasValidLocalization:
    trigger_emergency_land() # lost position estimate

```

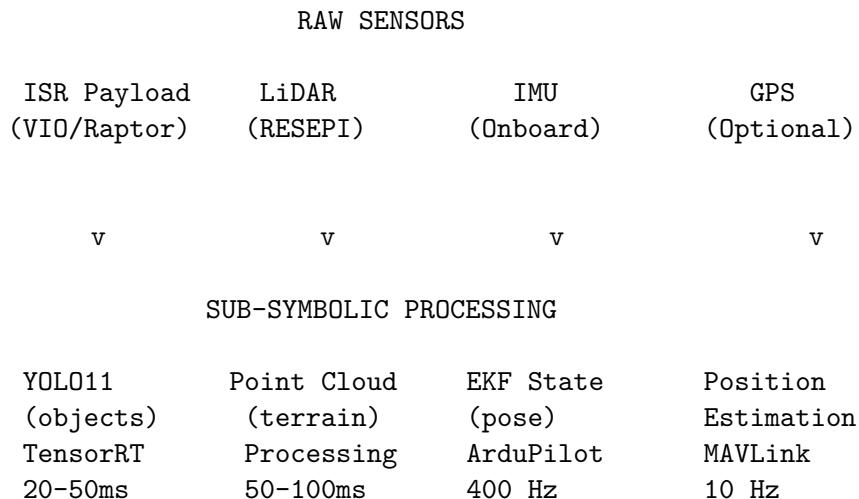
1.3.4.3 Action Filtering Workflow



1.3.5 5. Perception-to-Reasoning Bridge

Challenge: Vision models output sub-symbolic representations (bounding boxes, masks, embeddings), but reasoning engines need symbolic facts (relations, concepts, predicates) **Solution:** Symbolic abstraction layer that grounds perceptions to ontological concepts

1.3.5.1 Architecture: From Sensors to Symbols



v

SYMBOLIC ABSTRACTION LAYER (ROS 2 Grounding Nodes)

ObjectGroundingNode:

Input: DetectionArray (YOLO bboxes)
 Output: objectType(obj_123, 'person').
 position(obj_123, [x, y, z]).
 confidence(obj_123, 0.92).
 inRegion(obj_123, zone_A).

TerrainGroundingNode:

Input: Segmentation masks
 Output: terrainType(region_5, 'water').
 traversable(region_5, false).
 terrainSlope(region_5, 15.3).

SpatialRelationGroundingNode:

Input: Depth map + object positions
 Output: distance(drone, obj_123, 5.2).
 between(drone, obj_45, obj_67).
 northOf(obj_123, waypoint_A).
 clearance(forward, 8.5).

EventDetectionNode:

Input: Object tracking history
 Output: enters(obj_123, no_fly_zone).
 loitering(obj_45, duration(30)).
 violates(mission, constraint_7).

VLMGroundingNode (Optional):

Input: RGB frame + depth
 Prompt: "Describe objects and spatial relations as
 Prolog facts using SUMO ontology"
 Output: Structured facts for complex scene understanding

v

PROLOG KNOWLEDGE BASE (SWI-Prolog)

% Dynamic facts (asserted by perception grounding nodes)
 objectType(obj_123, person).

```

position(obj_123, [45.2, -122.1, 100]).  

distance(drone, obj_123, 5.2).  

terrainType(region_5, water).  

clearance(forward, 8.5).  
  

% Static rules (compiled from planning phase)  

safeToFlyOver(Region) :-  

    terrainType(Region, Type),  

    not(Type = water),  

    not(Type = urban).  
  

mustAvoid(Object) :-  

    objectType(Object, person),  

    distance(drone, Object, Dist),  

    Dist < 50. % meters  
  

hasSafeClearance(Direction) :-  

    clearance(Direction, Dist),  

    Dist > 3.0. % minimum safety margin  
  

violatesSafetyConstraint(Constraint) :-  

    mustAvoid(Object),  

    distance(drone, Object, Dist),  

    Dist < 10. % critical threshold

```

v

RL AGENTS + BEHAVIOR TREE EXECUTOR

Query examples:

- safeToFlyOver(current_region)? → bool
- mustAvoid(X)? → [obj_123, obj_45]
- violatesSafetyConstraint(C)? → abort/replan
- hasSafeClearance(forward)? → bool

Action selection constrained by Prolog query results

1.3.5.2 Why This Multi-Layer Approach?

Vision Models Alone Cannot: 1. **Reason relationally:** YOLO says “person at (x,y)”, not “person between drone and target” 2. **Check constraints:** Can’t encode “must avoid urban areas when communications-denied” 3. **Apply temporal logic:** Can’t reason “if in no-fly zone for >5 seconds, abort mission” 4. **Understand mission context:** Don’t know what “safe”, “compliant”, or “appropriate” means

The Symbolic Layer Provides: 1. **Semantic grounding:** Maps pixels → concepts from SUMO

ontology 2. **Spatial relations:** Computes n-ary relations (between, northOf, inside) 3. **Constraint evaluation:** Checks compiled mission rules against perceived world state 4. **Event recognition:** Detects complex events (entering zones, loitering, violations) 5. **Explainability:** Every decision traceable to symbolic facts and rules

1.3.5.3 Multi-Sensor Fusion Example

Scenario: Approaching person during autonomous navigation

EKF Pose:	Thermal/RGB:	YOLO11:
Velocity	Range at	Detection
[2.0, 0, 0]	bearing 0°:	class: person
m/s forward	5.2 meters	conf: 0.92
		bbox: [x,y,w,h]

v

Grounding Nodes (ROS 2)

```
ObjectGroundingNode:
objectType(obj_123,
           person).
position(obj_123,
         [5.2, 0, 0]).
```

```
SpatialRelationNode:
distance(drone,
          obj_123, 5.2).
timeToContact(obj_123,
              2.6).
```

v

Prolog KB Query

```
?- mustAvoid(X).
X = obj_123.

?- violatesSafetyConstraint
(C).
C = proximity_alert.
```

v

AGGM Triggered
Case 3: High-importance
safety concept

Backward reasoning:
→ Goal: reach safe distance
→ Priority: Safety (J) > B

v

RL Policy (Constrained)

Query Prolog for valid actions:
- Forward: INVALID (collision)
- Stop: VALID
- Avoid left: VALID
- Avoid right: VALID

Select: avoid(left, 5m)

v

Execute Action
[vx=0, vy=2.0, vz=0]

Monitor: distance(drone,
obj_123, Dist)
Until: Dist > 10

1.4 Implementation Roadmap

Detailed Plan: See [SYNTHESIS.qmd - Recommended Next Steps](#)

1.4.1 Phase 1: Ontology Foundation & Two-Phase Architecture (Weeks 1-6)

1.4.1.1 Part A: Planning Mode Infrastructure (Weeks 1-3)

- Create Podman container with SUMO ontology (SUO-KIF format), Vampire, Clingo, E-Prover
- All dependencies fully documented in `Containerfile.planning`
- Develop UAV domain ontology in SUMO (flight phases, spatial relations, sensors)
- Create Flyby-F11 application ontology (missions, constraints, NDAA compliance)
- Implement safety axioms in SUO-KIF, verify with Vampire

Deliverables: - /flyby-f11/ontology/Containerfile.planning (self-contained planning environment) - /flyby-f11/ontology/planning_mode/sumo_base.kif - /flyby-f11/ontology/planning_mode/uav_domain.kif - /flyby-f11/ontology/planning_mode/flyby_mission.kif - /flyby-f11/ontology/planning_mode/safety_axioms.kif - Verification scripts (Vampire proofs for safety properties)

1.4.1.2 Part B: Execution Mode Infrastructure (Weeks 4-6)

- Create Podman container with SWI-Prolog (ARM build for Jetson compatibility)
- All dependencies documented in Containerfile.execution
- Develop SUMO → Prolog translation tools (manual + semi-automatic)
- Compile critical safety axioms to Prolog rules
- Benchmark Prolog inference latency and memory footprint (both x86 dev + ARM Jetson)

Deliverables: - /flyby-f11/ontology/Containerfile.execution (self-contained execution environment) - /flyby-f11/ontology/execution_mode/compiled_rules.pl - /flyby-f11/scripts/sumo_to_prolog_translator.py - Translation validation tests (SUMO proofs Prolog queries) - Performance benchmarks (query latency <10ms, memory <100MB)

1.4.2 Phase 2: Perception-to-Reasoning Bridge (Weeks 7-10)

1.4.2.1 Part A: Symbolic Abstraction Layer (Weeks 7-8)

- Create Podman container with ROS 2 Humble + GPU passthrough
- All ROS 2 and perception dependencies in Containerfile.ros2
- Create ROS 2 grounding nodes package
- Implement ObjectGroundingNode (YOLO → Prolog facts)
- Implement TerrainGroundingNode (Segmentation → traversability facts)
- Implement SpatialRelationGroundingNode (depth + position → n-ary relations)
- Implement EventDetectionNode (tracking → temporal events)

Deliverables: - flyby_f11_ros2_ws/Containerfile.ros2 (ROS 2 + GPU passthrough) - flyby_f11_ros2_ws/Containerfile.vision (TensorRT + vision models) - flyby_f11_ros2_ws/src/perception - Unit tests for each grounding node - Integration tests (vision models → Prolog assertions)

1.4.2.2 Part B: Phase Transition Manager (Weeks 9-10)

- Implement mission planner node (runs heavyweight reasoners in planning container)
- Implement phase transition controller (container orchestration: planning → execution)
- Develop memory profiling and monitoring tools
- Test full planning → execution workflow with container switching

Deliverables: - flyby_f11_ros2_ws/src/mission_planner/ - flyby_f11_ros2_ws/src/phase_transition_manager/ - podman-compose.yml (orchestrates planning + execution containers) - Memory allocation tests (16GB budget validation) - Latency benchmarks (phase transition time)

1.4.3 Phase 3: Multi-Agent RL (Weeks 9-14)

- Define MDPs for mission planner, behavior selector, trajectory optimizer
- Implement AGGM (6-stage process) for each agent
- Develop training infrastructure (Gymnasium environments, reward shaping)

- Train policies with experience sharing (NPS cluster)

Deliverables: - `flyby_f11_ros2_ws/src/ontology_r1/` - Trained policies (mission/behavior/trajectory agents) - Ablation study results (with/without ontology constraints)

1.4.4 Phase 4: Benchmark Evaluation (Weeks 15-18)

- Develop scenarios (waypoint nav, obstacle avoid, GPS-denied, mission adapt)
- Implement baselines (pure RL, pure rules, LLM-based AOA)
- Collect metrics (SR, OSR, DTS, SPL, collision rate, efficiency)
- Statistical analysis (significance testing, Pareto frontiers)

Deliverables: - Benchmark suite with evaluation scripts - Comparative results (tables, plots, failure analysis) - Technical report / conference paper draft

1.4.5 Phase 5: Hardware Validation (Weeks 19-24)

- **Project-drone testing** (Weeks 19-21): Indoor waypoint nav, obstacle avoid, GPS-denied
- **Flyby-F11 integration** (Weeks 22-23): Outdoor survey/inspection missions
- **Stress testing** (Week 24): Sensor failures, environmental challenges, edge cases
- Iterative refinement based on flight logs

Deliverables: - Flight test videos with ontology decision overlays - Hardware performance benchmarks (Jetson utilization) - Safety assessment report (collision-free hours)

1.4.6 Phase 6: Documentation (Weeks 25-26)

- System architecture documentation
- Research paper for ICRA/IROS/RSS
- Open-source repository release
- MCTSSA demonstration package

1.5 Key Innovations and Contributions

1.5.1 1. Two-Phase Compute Architecture for Edge Ontological Reasoning

Challenge: Heavyweight ontological reasoners (Vampire, full SUMO) require >8GB RAM and multi-second inference times, incompatible with real-time UAV control and concurrent vision processing on 16GB unified memory **Solution:** Novel two-phase architecture separating planning and execution - **Planning Mode:** 100% compute for heavyweight reasoning (SUMO + Vampire/Clingo) during mission receipt/replanning - **Execution Mode:** Compiled Prolog rules (~100MB) coexist with vision models (~12GB) for real-time inference - **Phase Transition:** Automatic model swapping via memory manager (unload reasoners → load YOLO/segmentation) - **Performance:** <10ms Prolog query latency, supports concurrent 10Hz reasoning + 20Hz vision processing

Expected Impact: - First demonstration of full SUMO ontology verification + real-time execution on edge hardware - Enables rigorous safety proofs (Vampire) without sacrificing runtime performance - Memory-constrained platforms can now run both complex reasoning and modern vision models

1.5.2 2. Perception-to-Reasoning Bridge for Symbolic Grounding

Challenge: Vision models output sub-symbolic representations (bounding boxes, embeddings) while ontological reasoners require symbolic facts (predicates, relations) **Solution:** ROS 2 grounding nodes that translate perceptions to ontology-aligned Prolog facts - ObjectGroundingNode: YOLO detections → `objectType(obj_123, person)` - SpatialRelationGroundingNode: Depth maps → n-ary relations `between(drone, obj1, obj2)` - EventDetectionNode: Tracking history → temporal predicates `enters(obj, zone)` - VLMGroundingNode (optional): Scene understanding → structured symbolic facts

Expected Impact: - Closes the “semantic gap” between perception and reasoning - Enables SUMO’s rich spatial reasoning (ternary+ relations) from sensor data - Provides explainability: every decision traces through symbolic facts to axioms

1.5.3 3. Real-Time Ontological Reasoning on Edge Hardware

Challenge: Traditional ontology reasoners require server-class compute **Solution:** Hybrid SUMO (planning) + SWI-Prolog (execution) approach - SUMO: Full expressivity for verification (n-ary relations, FOL) - Prolog: Lightweight runtime (<100MB, <10ms queries) - Translation validation: Equivalence testing between SUMO proofs and Prolog - Target: 10 Hz reasoning loop concurrent with 20 Hz vision processing

Expected Impact: First demonstration of SUMO-grade reasoning on ARM edge platform for UAV control

1.5.4 4. AGGM for GPS-Denied Navigation

Challenge: Pure RL fails in unseen situations (UAV-ON: 7.30% success rate) **Solution:** Backward reasoning to create goals returning to known safe states - Ontology defines state similarity metrics - Priority function balances safety (J) vs. task completion (B) - Runtime goal generation without retraining

Expected Impact: >50% success rate in novel environments (vs. 7.30% baseline), <10% collision rate (vs. 37-65% baseline)

1.5.5 5. Explainable Autonomy for Defense Applications

Challenge: Black-box RL policies unacceptable for MCTSSA collaboration **Solution:** Semantic decision trace through ontological reasoning - Every action justified by SWRL rule firing - Safety constraints auditable and verifiable - Mission-intent interpretation transparent to operators

Expected Impact: First explainable autonomous UAV system meeting NDAA compliance and defense safety standards

1.5.6 6. Hierarchical Multi-Agent RL with Ontological Structure

Challenge: Monolithic policies don’t scale to complex missions **Solution:** Specialized agents at each hierarchy level (mission/behavior/trajectory) - Experience sharing via shared replay buffer - Policy distillation from expert to novice agents - Meta-learning across mission types

Expected Impact: Improved sample efficiency (faster training convergence), modular extensibility (add new missions/behaviors without full retraining)

1.6 Success Criteria

1.6.1 Simulation Benchmarks (Phase 4)

- Success Rate (SR) > 50% in novel environments (UAV-ON-style)
- Collision Rate < 10% (vs. 37-65% for unconstrained baselines)
- Safety Constraint Violations = 0 (hard requirement)
- Explainability: 100% of actions traceable to ontology rules

1.6.2 Hardware Performance (Phase 5)

- Ontology reasoning latency < 100ms at 10 Hz
- Jetson Orin NX resource utilization < 70% (CPU/GPU/memory)
- Real-time control loop: 10 Hz for mission, 100 Hz for trajectory
- Successful GPS-denied navigation for >5 minutes using EKF-only localization

1.6.3 Flight Testing (Phase 5)

- 10+ successful autonomous missions (waypoint navigation)
- 5+ successful obstacle avoidance scenarios (dynamic obstacles)
- 3+ successful mission adaptations (unexpected events)
- Zero collisions in controlled testing environment
- MCTSSA demonstration: communications-denied mission completion

1.6.4 Research Impact (Phase 6)

- Conference paper accepted (ICRA/IROS/RSS)
 - Open-source release with 100+ GitHub stars
 - Contribution to UAV-ON benchmark (ontology-constrained baseline)
 - MCTSSA collaboration continuation (follow-on projects)
-

1.7 References

Detailed Literature Review: [SYNTHESIS.qmd](#)

Core Papers: 1. Hare & Tang (2024) - Multi-agent ontology-driven RL for personalized systems 2. Ghanadbashi & Golpayegani (2022) - AGGM for unseen situations (traffic control) 3. Aguado et al. (2024) - Survey of ontology-enabled robot dependability 4. UAV Collision Avoidance Ontologies - Domain application review 5. Xiao et al. (2025) - UAV-ON benchmark for object-goal navigation 6. Hu et al. (2025) - Survey of hybrid decision-making for autonomous vehicles

Platform Details: [SYSTEM_CONSTRAINTS.qmd](#)

Ontology Specification: [ONTOLOGY_FOUNDATION.qmd](#)

1.8 Appendix: Alignment with Project-Drone

Development Strategy: Shared autonomy components developed on accessible hardware

project-drone (Development Platform)

Hardware: Jetson Orin Nano Super 8GB (67 TOPS)

Sensors: T265 visual odometry, D455 depth camera

Purpose: Algorithm development, rapid prototyping

Shared packages:

autonomy_core/ (mission planning, waypoint navigation)

behavior_trees/ (BT mission logic, ontology-aware nodes)

perception_pipeline/ (vision models, semantic fusion)

px4_interface/ (MAVSDK bridge, flight abstraction)

↓ (symlinks when ready)

flyby-f11 (Deployment Platform)

Hardware: Jetson Orin NX 16GB (50 TOPS, 2x memory)

ISR Payloads: Gremsy VIO | RESEPI LiDAR | NextVision Raptor

Purpose: Mission-specific deployment (MCTSSA)

Packages:

flyby_f11_bringup/ (mission launch configurations)

flyby_f11_sensors/ (platform-specific drivers)

flyby_f11_mission/ (MCTSSA mission logic)

ontology_interface/ (reasoner, perception mapper, action validator)

Benefit: Ontology-constrained RL developed incrementally on project-drone, then deployed to flyby-f11 when hardware available. Shared packages ensure consistency while platform-specific packages handle hardware differences.

Document Version: 2.0 **Last Updated:** 2024-12-26 **Status:** Architecture updated based on Phase 3 evaluation. Phases 1-3 complete, Phase 4+ updated for single-reasoner (Vampire) architecture.

Changelog: - v2.0 (2024-12-26): Updated to single-reasoner architecture (Vampire only) based on Phase 3 evaluation - v1.0 (2024-12-25): Initial architecture with two-phase compute (SUMO+Vampire planning, SWI-Prolog execution)